



# **The Autonomous, Understandable, Research Assistant (A.U.R.A.):** Innovating Customer Insights

“A self-reflective RAG framework  
with explainable governance for  
research tasks.”



# The Problem: The Research Bottleneck

- Up to 80% of a data professional's time goes into data preparation, gathering and organizing instead of analysing  
*(Forbes / CrowdFlower survey, 2016).*
- The growing volume of available data is making the analyst's job increasingly complex and time-consuming.
- Clients need a tool that delivers fast, consistent, and relevant research data to address the challenges their analysts face.



**The Goal is to improve:**  
Research Efficiency • Consistency • Transparency

# Using Responsible AI Guardrails

- Accountability & Oversight
- Transparency & Explainability
- Empowerment, not Replacement
- Security & Bias Mitigation
- Data is used responsibly and within reason

"AI should be used to do the heavy lifting, but important decisions should still be made by people with years of experience and technical expertise."

# Understanding the right approach: Reliability Over Pure Generation

|                       | Generative AI (Purely Generative)                         | Retrieval-Augmented Generation (RAG-LLM)                                                                        |
|-----------------------|-----------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| <b>Functionality</b>  | Creates new text purely from its training data.           | Retrieves factual information from verified sources before generating text.                                     |
| <b>Usability</b>      | Very easy — users can prompt it directly.                 | Less intuitive for non-technical users — requires embedding setup, vector indexing, and retrieval configuration |
| <b>Strengths</b>      | Excellent for conversational, or summarisation tasks.     | Great for retrieving factual, auditable, and explainable research.                                              |
| <b>Data Creation</b>  | Can “hallucinate” or produce inaccurate/outdated content. | Data often unstructured, requiring extra cleaning and formatting before analysis                                |
| <b>Data Integrity</b> | Black-box outputs with limited explainability.            | Every insight links back to a real document making data transparent and traceable                               |

# My proposal: A.U.R.A.

## A three layer solution

### Query Layer



### Retrieval Layer



### Governance Layer

#### Understanding Intent

(LLM Router)

Transforms analyst questions into structured, context-aware search plans.

- Interprets intent, entities, and scope using a language model.
- Chooses the best route — internal data, web sources, or fallback generation.

#### Retrieving Insights

(Vector + Web Search + CMD-R)

Retrieves and assembles evidence-based information from relevant sources.

- Combines vector search, web retrieval, and CMD-R generation.
- Produces grounded, cited drafts for review and refinement.

#### Validation and Assurance


(Critique + Governance Nodes)

Evaluates and governs the accuracy and completeness of retrieved insights.

- Scores confidence and detects hallucinations or missing context.
- Refines or re-queries when necessary to deliver verified, self-governed results.

# How the AURA Operates Conceptually

## *LangGraph RAG Framework*




"Built using LangGraph for orchestration of these three layers."

# From Concept to Technical Implementation

| Concept Idea (Ask<br>→ Find → Verify)                                                                                                                     | LangGraph Components                                                          | Output & Responsibility                                             |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|---------------------------------------------------------------------|
| Query Layer “Ask”                                                                                                                                         | Query Router Node (LLM) – classifies intent & routes to retrievers            | Clarified, structured query with retrieval plan                     |
| Retrieval Layer “Find”                                                                                                                                    | Retriever Nodes (Vector DB + Web Search) → CMD-R Generator                    | Grounded draft with citations from both internal & external sources |
| Governance Layer “Verify”                                                                                                                                 | Critique Node + Governance Node within a Repair Loop (LangGraph control flow) | Confidence score, bias report, and final verified summary           |
| Each layer is represented by LangGraph nodes that interact through a self-reflective feedback loop to ensure factual, transparent, and auditable outputs. |                                                                               |                                                                     |

# A.U.R.A. Self-Reflective RAG Flow (LangGraph Framework)



🧠 Query → 📖 Retrieve → 🛠️ Generate → 🔍 Critique + Govern → ♻️ Repair / ✅ Result



# Efficiency - Automating research flow for faster insight generation.

- ❑ Reduces time spent on manual research, allowing analysts to focus on higher-value interpretation and decision-making.
- ❑ Enables faster project turnaround as relevant sources and summaries can be rapidly retrieved and prepared for analysis. Extremely valuable in fast-moving markets
- ❑ Lets analysts perform more targeted research through the interactive query layer

**Before: 80% prep / 20% analysis → After: 20% prep / 80% analysis.**

# Consistency: Maintaining traceable, repeatable analytics.

- ❑ Layered design improves visibility, adaptability, and auditability across each stage
- ❑ Generates **citations and confidence scores** for every finding.
  - ❑ Ensures all insights remain **traceable and verifiable**.
- ❑ Delivers a structured, repeatable research process for analysts.

# Quality & Governance: Ensuring human-in-the-loop verification for responsible AI.

- ❑ Maintains human validation and judgement as a core part of the research process.
- ❑ Applies **BDO's GUARD framework** to ensure responsible AI governance and safeguarding of client information.
- ❑ Flags potential bias or misinformation through bias scoring on each sourced data point.

## **Next Step: Agentic Collaboration and Shared Intelligence**

Building on A.U.R.A.'s self-reflective RAG architecture.

“Expanding reflection into multi-agent collaboration and domain-specific fine-tuning, enabling contextual learning across departments.”

# Next Step: Agentic Collaboration and Shared Intelligence

Building on A.U.R.A.'s self-reflective RAG architecture.

## Shared Memory

Stores past queries, trusted sources, confidence scores, and critique feedback

## Feedback Loop (Self Improvement)

2-way communication

2-way communication

### Query Agent (Router Node)


Refines analyst intent & routing between internal data, web, or fallback generation. Learns from governance feedback to improve query formulation.

**Query Router (LLM)**  
decides between internal data, web search, or fallback generation

Question Input

### Retrieval Agent (Vector + Web Layer)

Sources and ranks factual data from vector and web searches. Adjusts scope using confidence and critique feedback.



### Summary/Governance Agent (Critique + Govern Nodes)

Validates relevance, evidence, and completeness. If confidence is low, triggers re-query or reframing instructions.

**Critique Agent**  
(Hallucination / Evidence real?)



Yes

**Governance Agent**  
(Does the response answer the question?)



Yes

Self-Governed Response

Agents wrap A.U.R.A.'s self-reflective RAG nodes, sharing memory and critique feedback to plan smarter re-queries, reduce loops, and coordinate insights across domains.